



## Understanding & Planning Hadoop Cluster



# Contents

1

What is Node?

2

Hard disk of a Hadoop Node

3

What is a DC, Cluster?

4

Hadoop / CDH Versions

5

Hadoop Architecture

6

HDFS Workflow

7

HDFS Write

8

HDFS Read

9

Hadoop General Planning Considerations

# What is a Node?



Foreground processes/ Appls

Background processes / Daemons

Operating System

Hardware

# Hard Disk of a Hadoop Node



# Understanding Hadoop Cluster



Cluster DC-East



DC-West



# Hadoop / CDH Versions

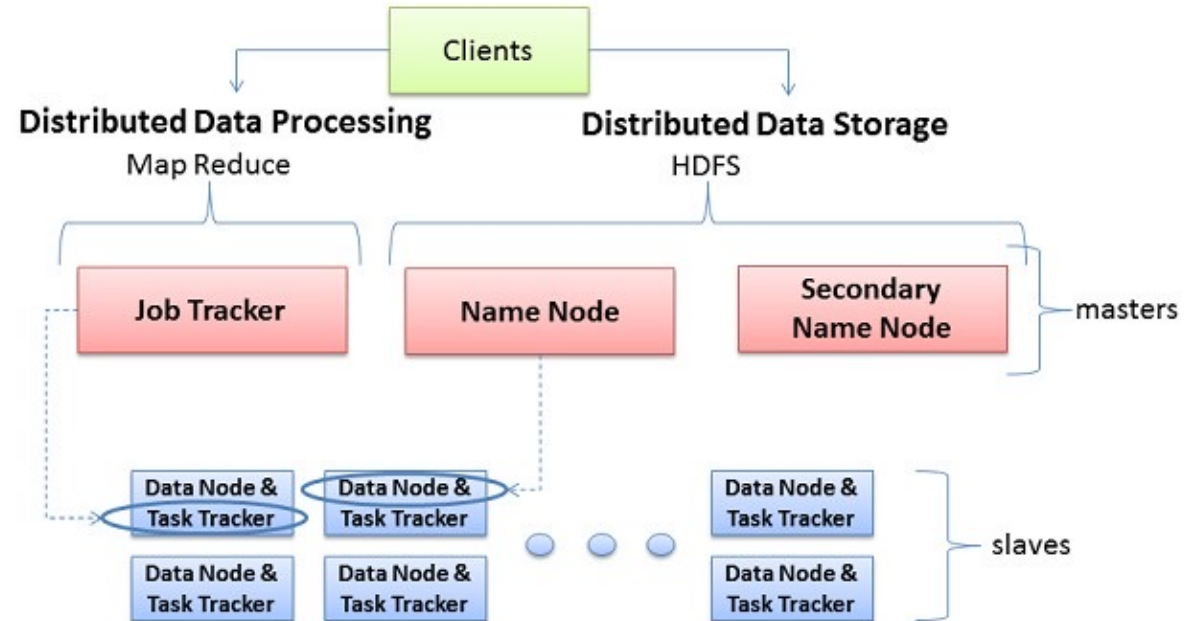


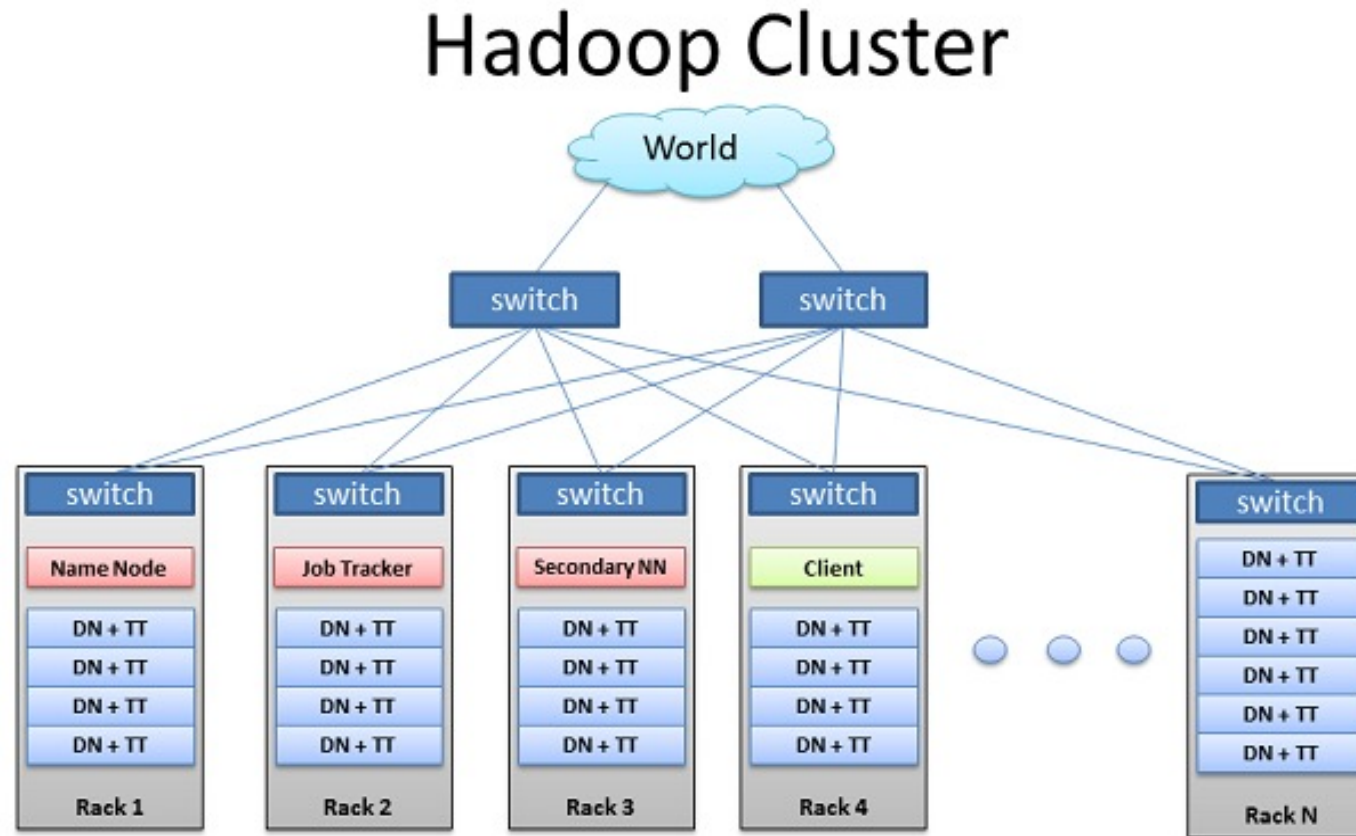
Hadoop	Cloudera(CDH)
0.20	<CDH 4.0
2.0	CDH 4.x
2.3	CDH5.0
2.6	CDH5.4,5.5,5.7
2.7	CDH 5.16
3.0	>CDH 6.0

## 3 Major Categories

- ☁ Clients
- ☁ Master Nodes
- ☁ Slave Nodes

## Hadoop Server Roles





This is the typical architecture of a Hadoop cluster



# Understanding Hadoop Cluster



Property File	Property Name/Purpose
hdfs-default.xml	configurations for HDFS <b>Ex:</b> dfs.datanode.data.dir, dfs.namenode.name.dir, dfs.replication
mapred-default.xml	configuration for MR <b>Ex:</b> mapreduce.jobtracker.http.address
core-default.xml	configuration across the cluster <b>Ex:</b> hadoop.tmp.dir, fs.default.name
hdfs-site.xml	Override storage(HDFS) specific properties
mapred-site.xml	Override any processing(MR) specific properties
core-site.xml	Override properties other than above 2



## Typical Workflow

- Load data into the cluster (HDFS writes)
- Analyze the data (Map Reduce)
- Store results in the cluster (HDFS writes)
- Read the results from the cluster (HDFS reads)

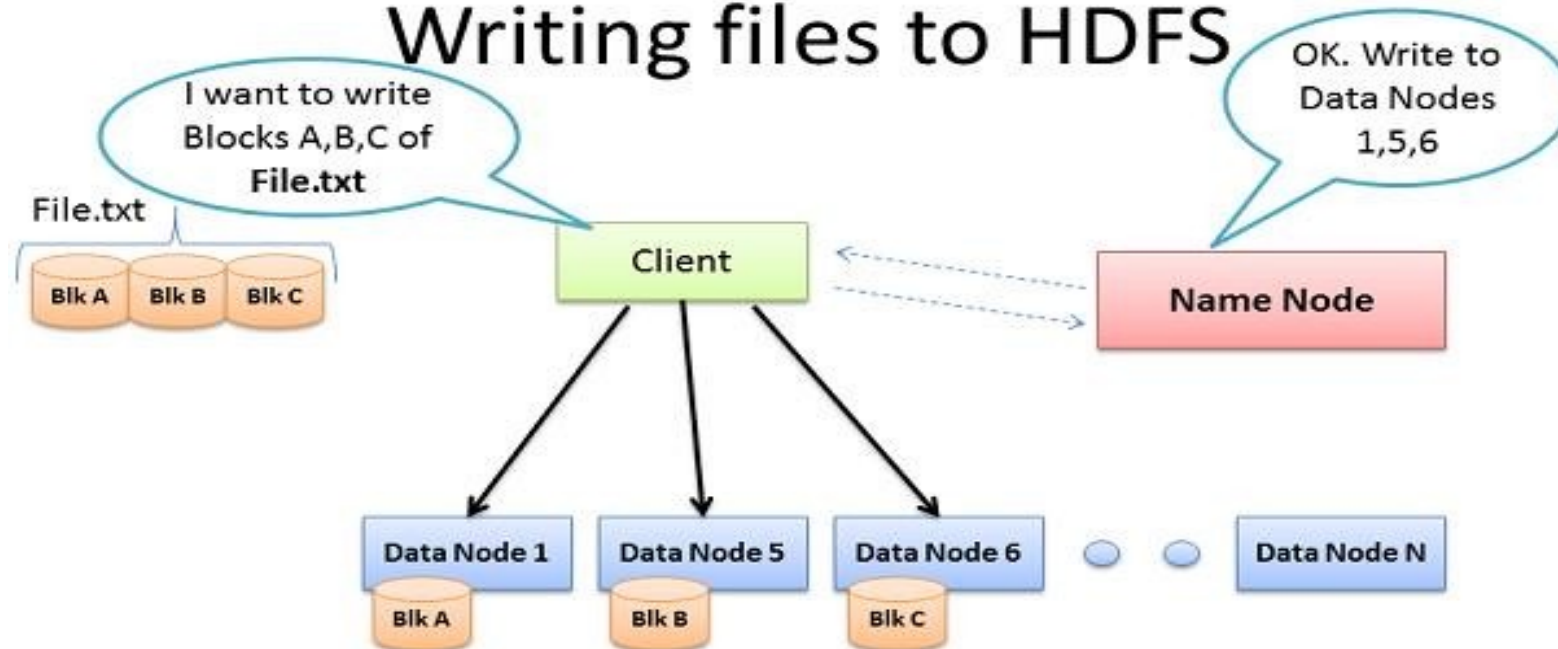
Sample Scenario:

How many times did our customers type the word **"Refund"** into emails sent to customer service?

Huge file containing all emails sent  
to customer service

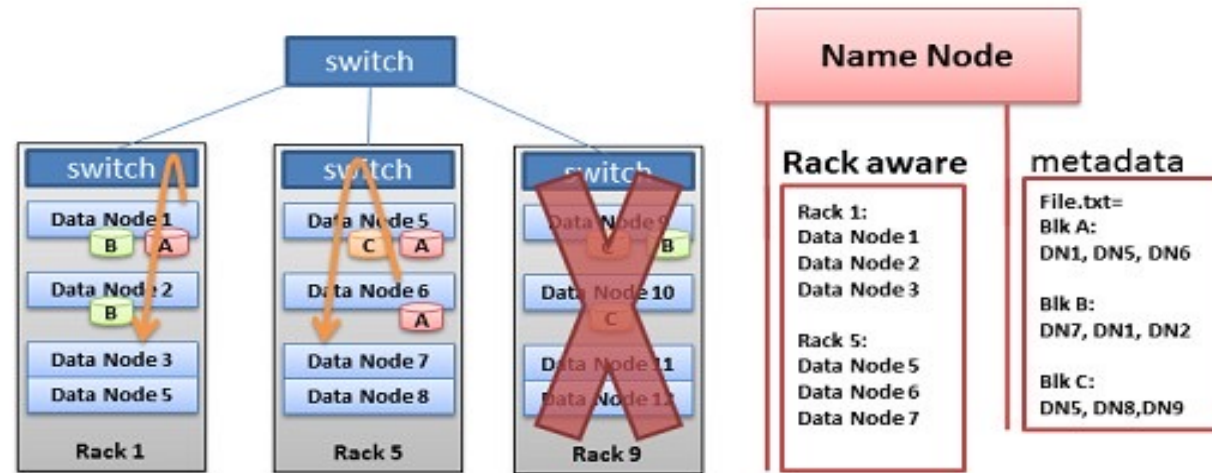


## Writing files to HDFS



- Client consults Name Node
- Client writes block directly to one Data Node
- Data Nodes replicates block
- Cycle repeats for next block

## Hadoop Rack Awareness – Why?



- Never loose all data if entire rack fails
- Keep bulky flows in-rack when possible
- Assumption that in-rack is higher bandwidth, lower latency

# Understanding Hadoop Cluster



## Name node replica placement Strategy

1st Replica → 1 Node (Usually client Node) say DN1 of Rack 1.

2nd Replica → Another Node Say, DN5 of another Rack 2.

3rd Replica → Same Rack R2 another Node say DN6.

## **By Rule of thumb**

1/3rd of Replicas in 1 Rack.

2/3rd of Remaining Replicas in Rack 2.

Remaining replicas are distributed across other racks equally.

**Ex:** Replication factor =9

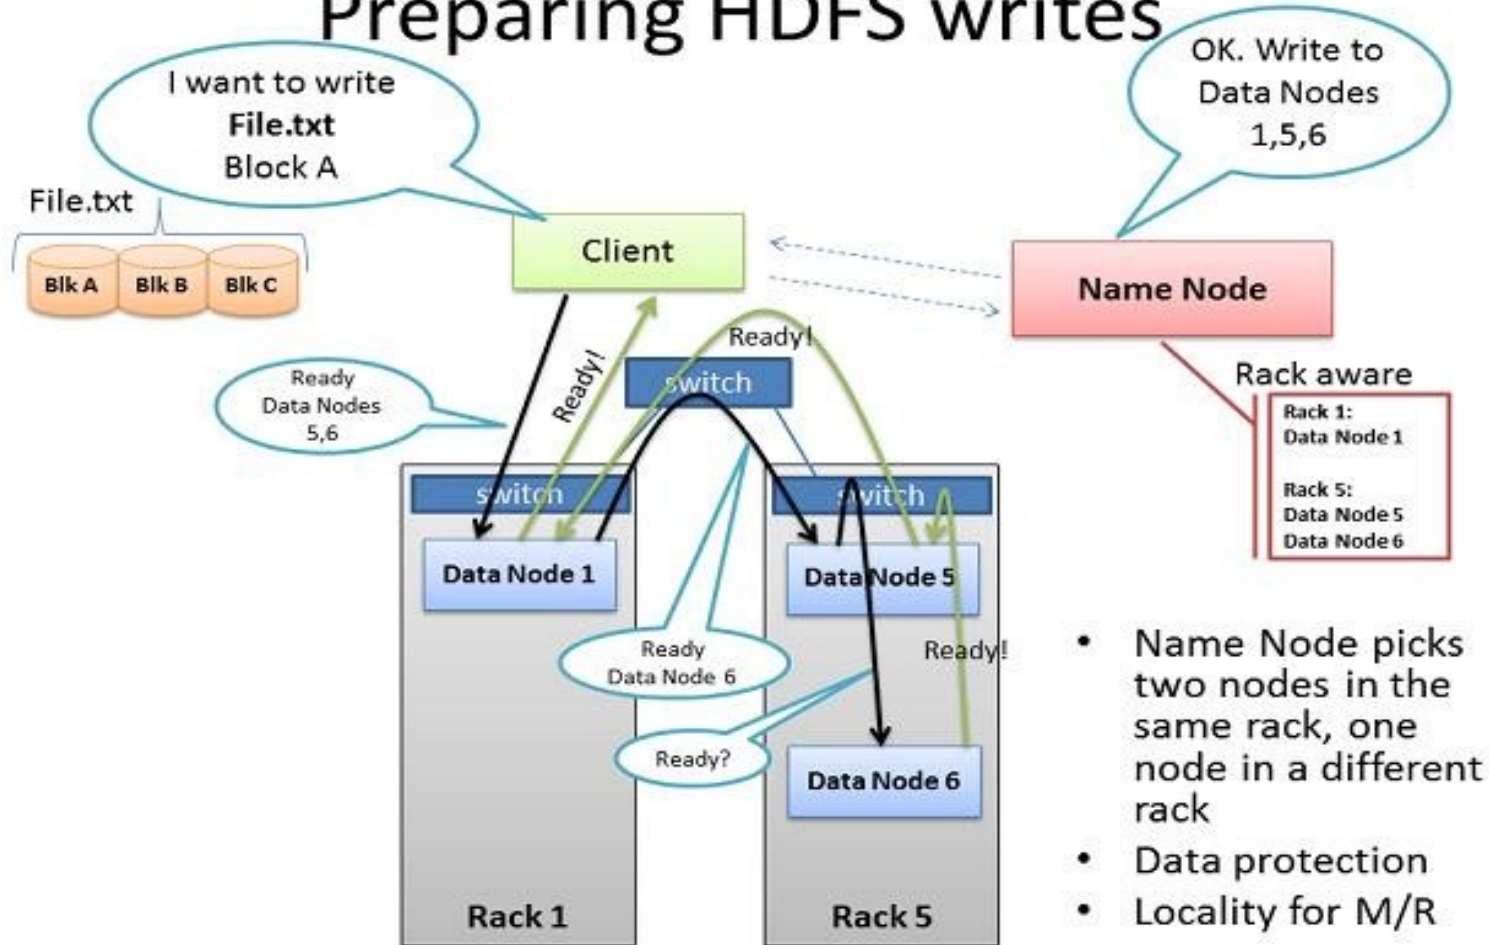
1/3rd = 3 replicas in 1 Rack.

2/3rd =  $2/3 * 6 = 4$  in another Rack.

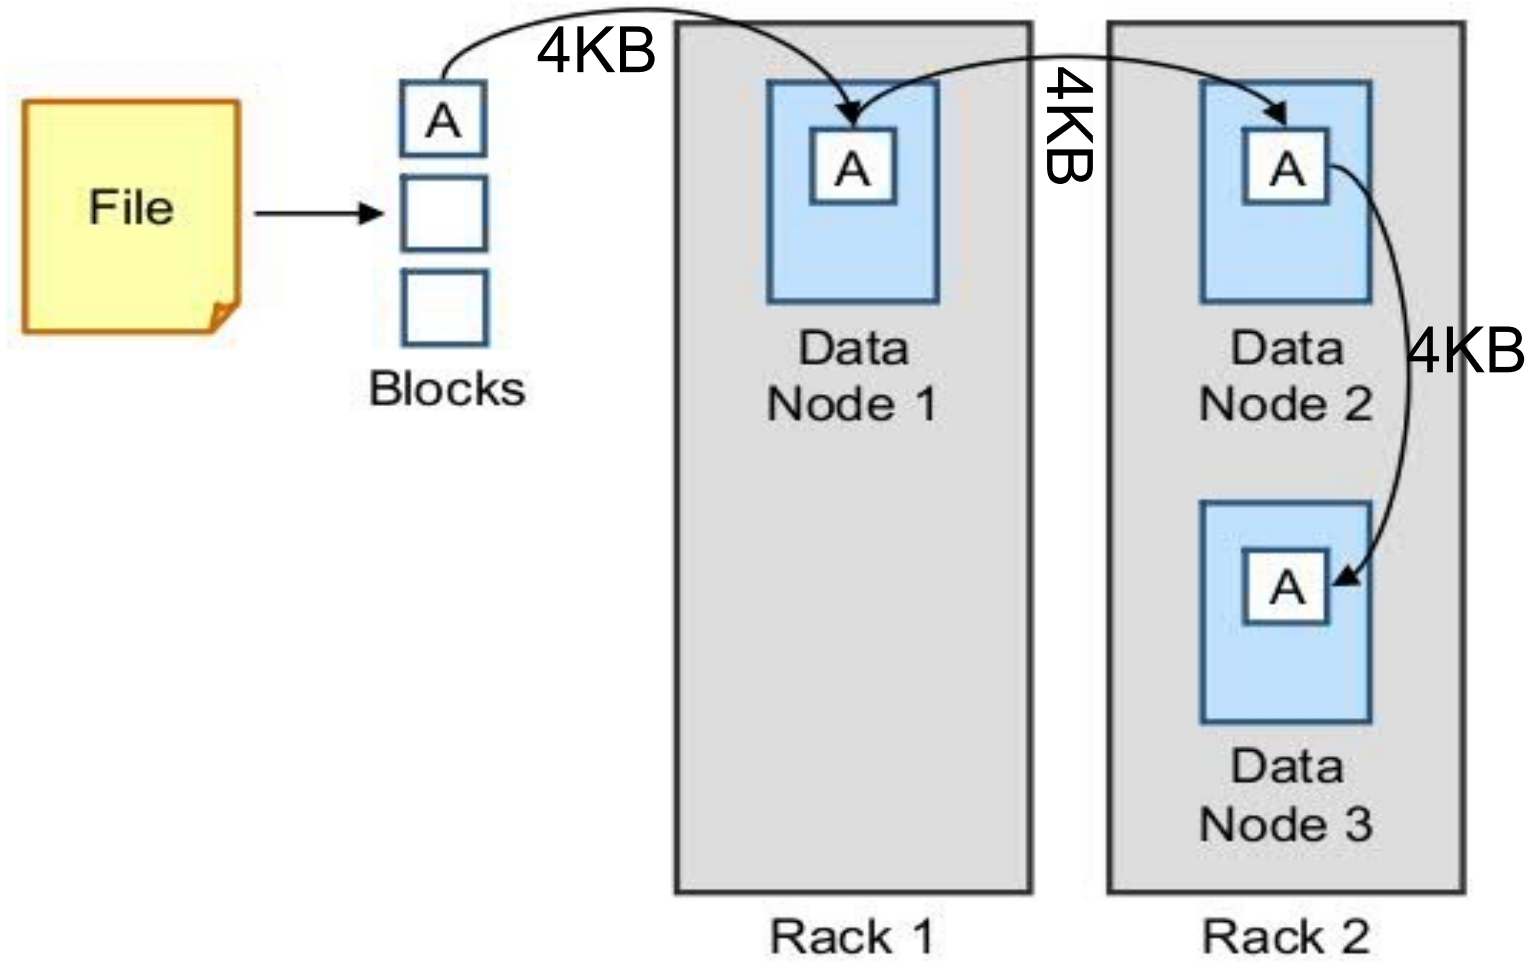
Remaining 2 replicas in other racks.

---

## Preparing HDFS writes



# Understanding Hadoop Cluster





## Name node Meta Data

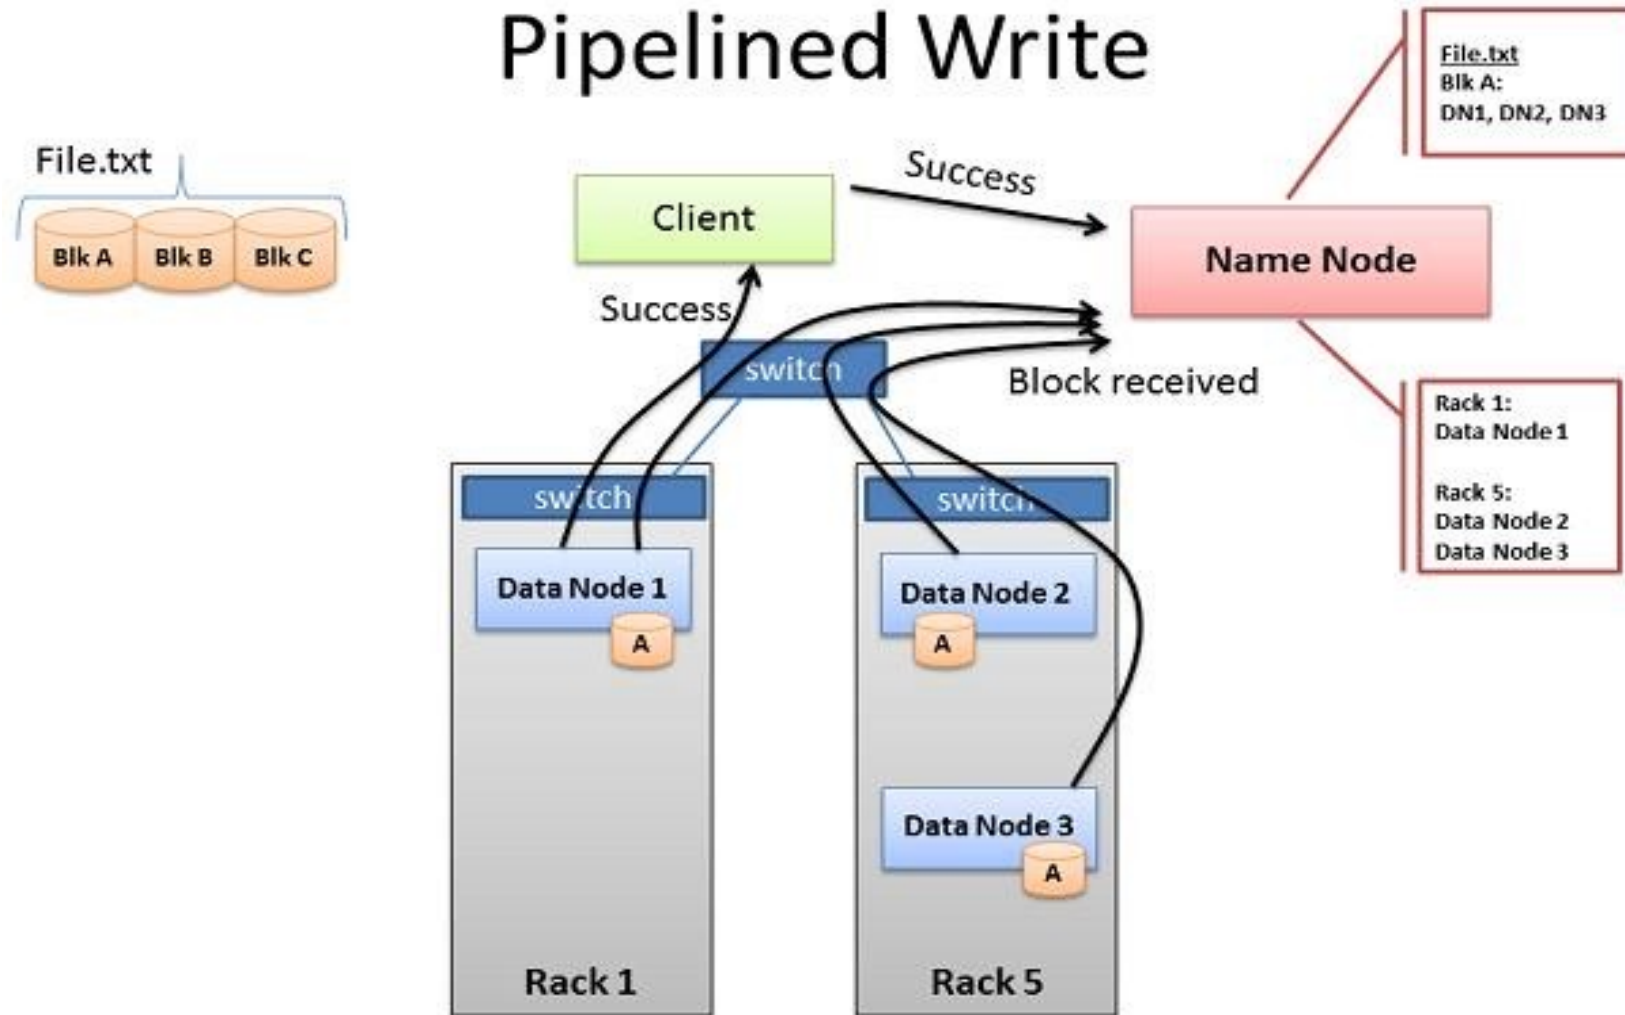
1. File to Block Mapping
  - Ex: File1.txt → Block A, Block B, Block C
2. Block To Node Mapping
  - Ex: Block A → Node1, Node 5, Node 6
  - Block B → Node 7, Node 1, Node 2
  - Block C → Node3, Node 8, Node 9



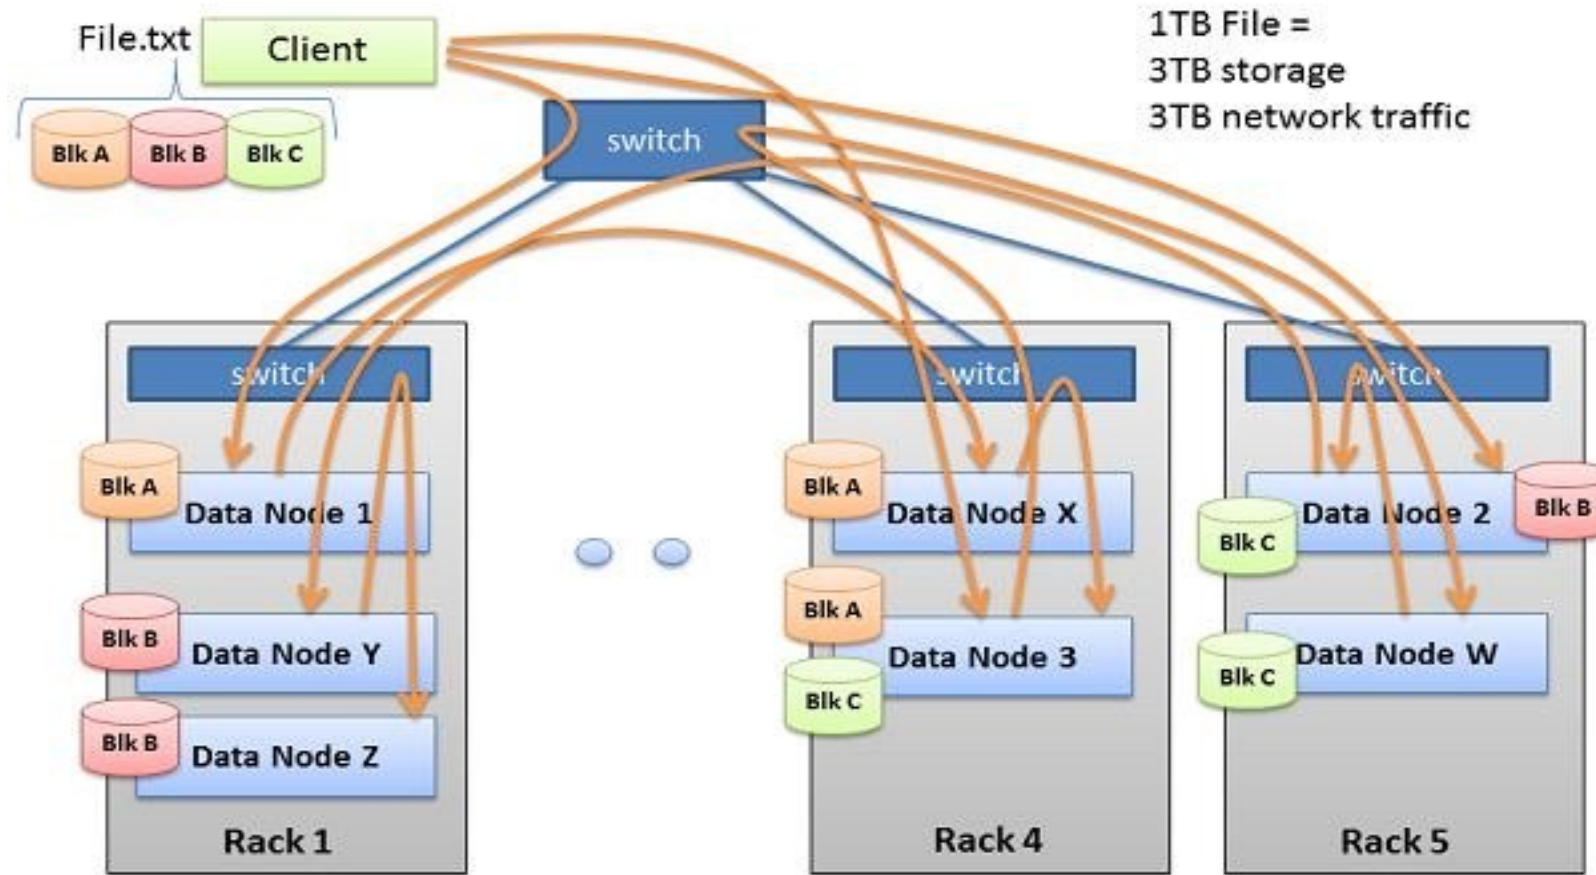
# Understanding Hadoop Cluster



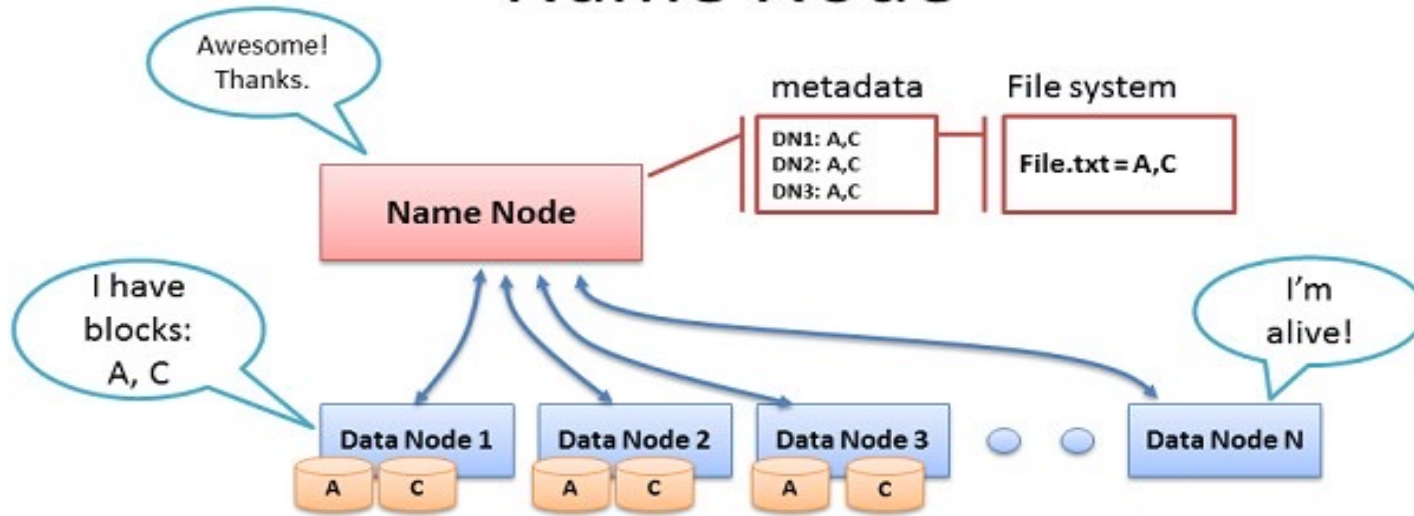
## Pipelined Write



## Multi-block Replication Pipeline

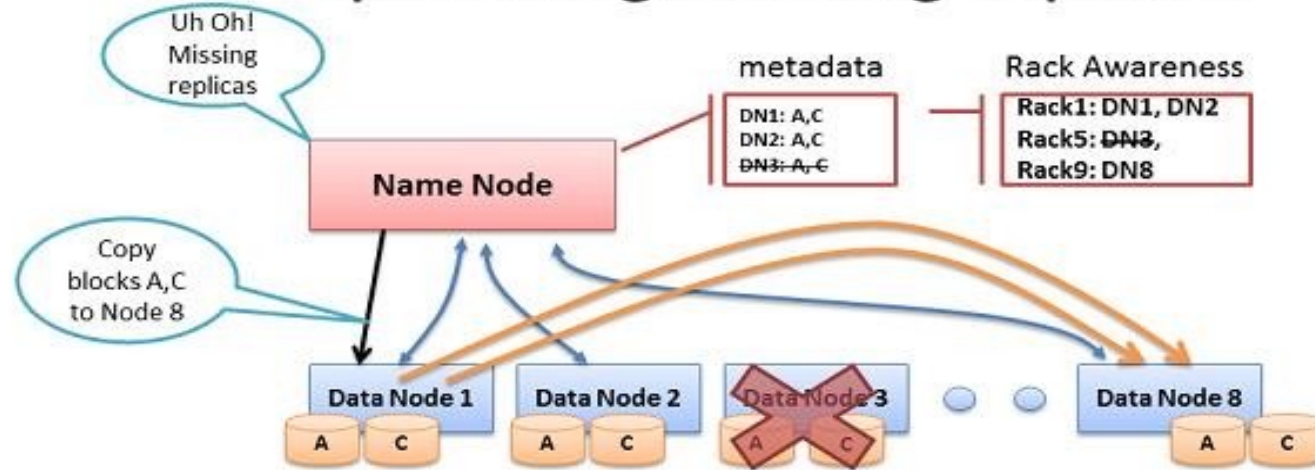


## Name Node



- Data Node sends Heartbeats
- Every 10<sup>th</sup> heartbeat is a Block report
- Name Node builds metadata from Block reports
- TCP – every 3 seconds
- If Name Node is down, HDFS is down

## Re-replicating missing replicas

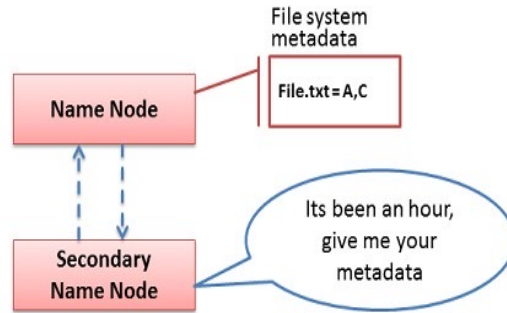


- Missing Heartbeats signify lost Nodes
- Name Node consults metadata, finds affected data
- Name Node consults Rack Awareness script
- Name Node tells a Data Node to re-replicate

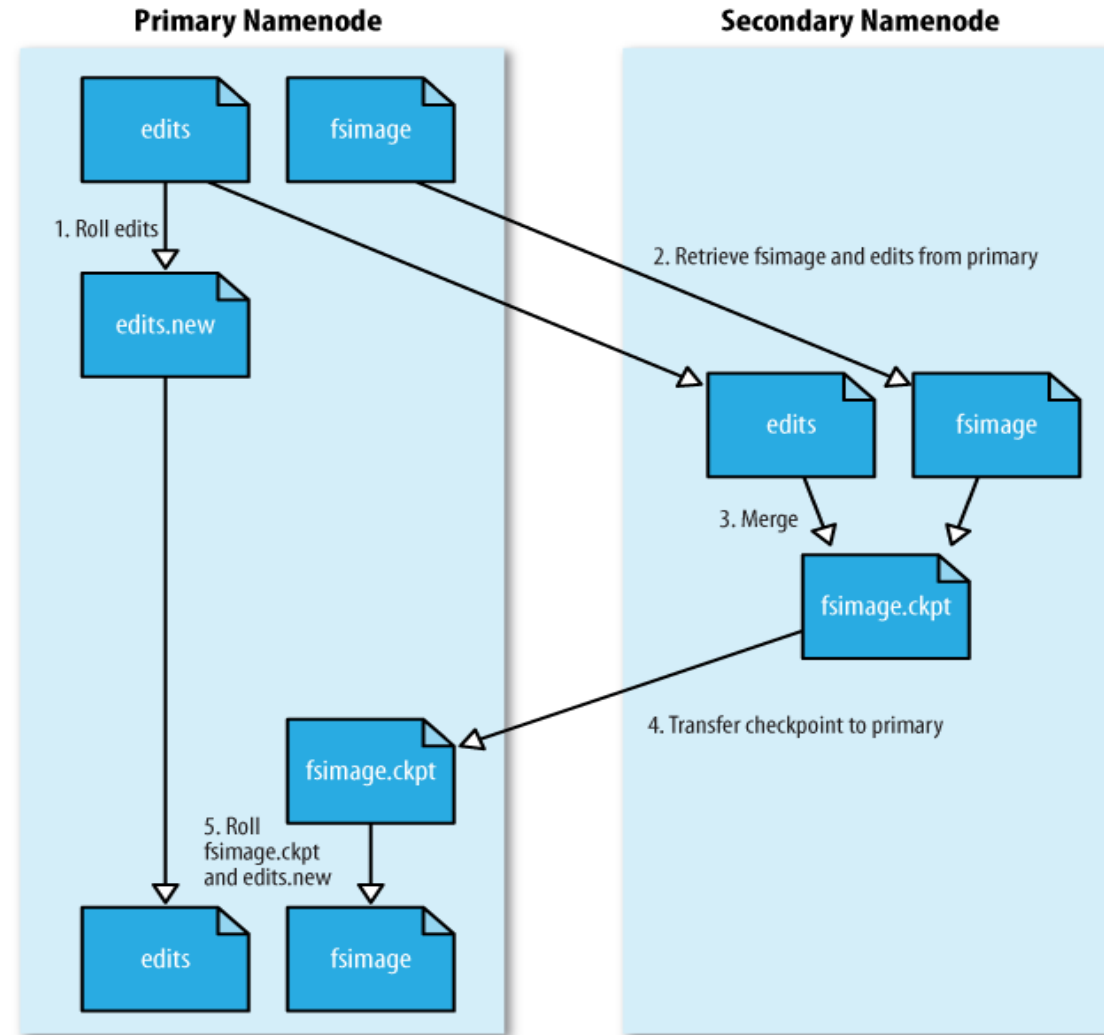
# Understanding Hadoop Cluster



## Secondary Name Node

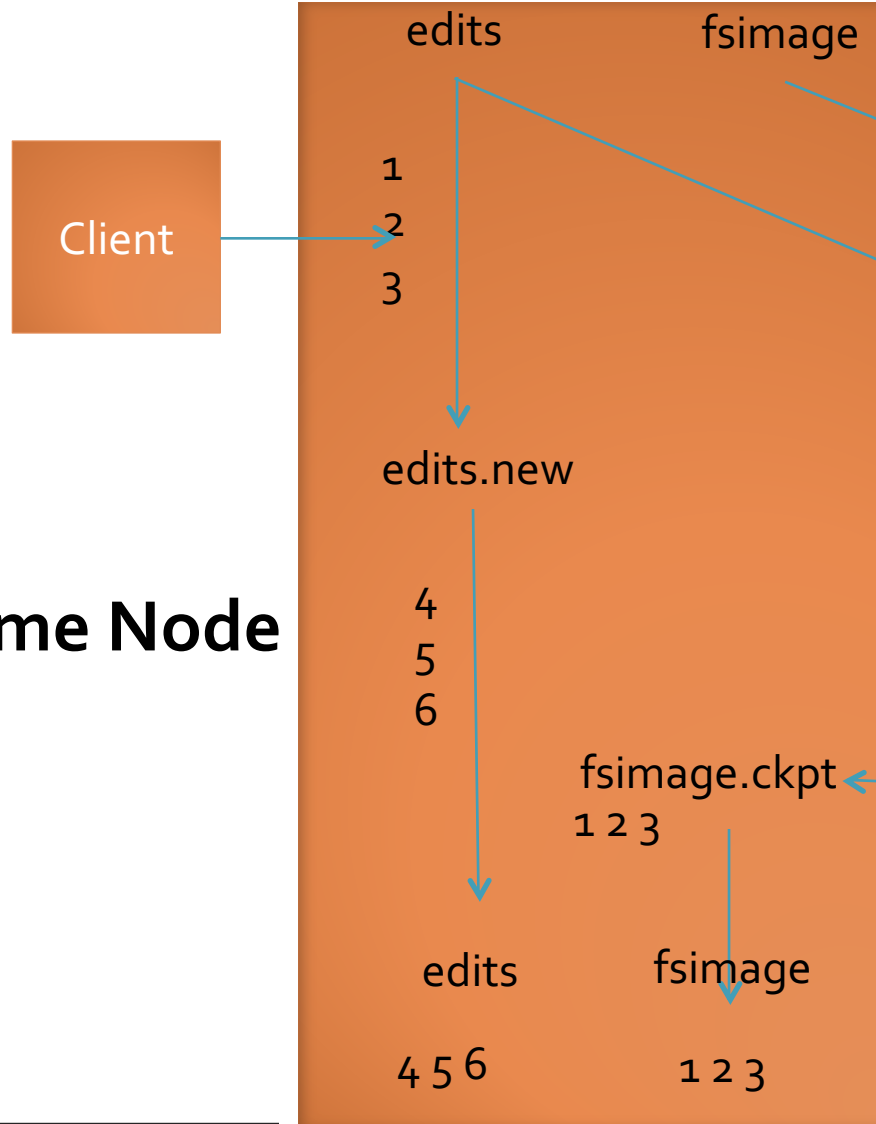


- Not a hot standby for the Name Node
- Connects to Name Node every hour\*
- Housekeeping, backup of Name Node metadata
- Saved metadata can rebuild a failed Name Node

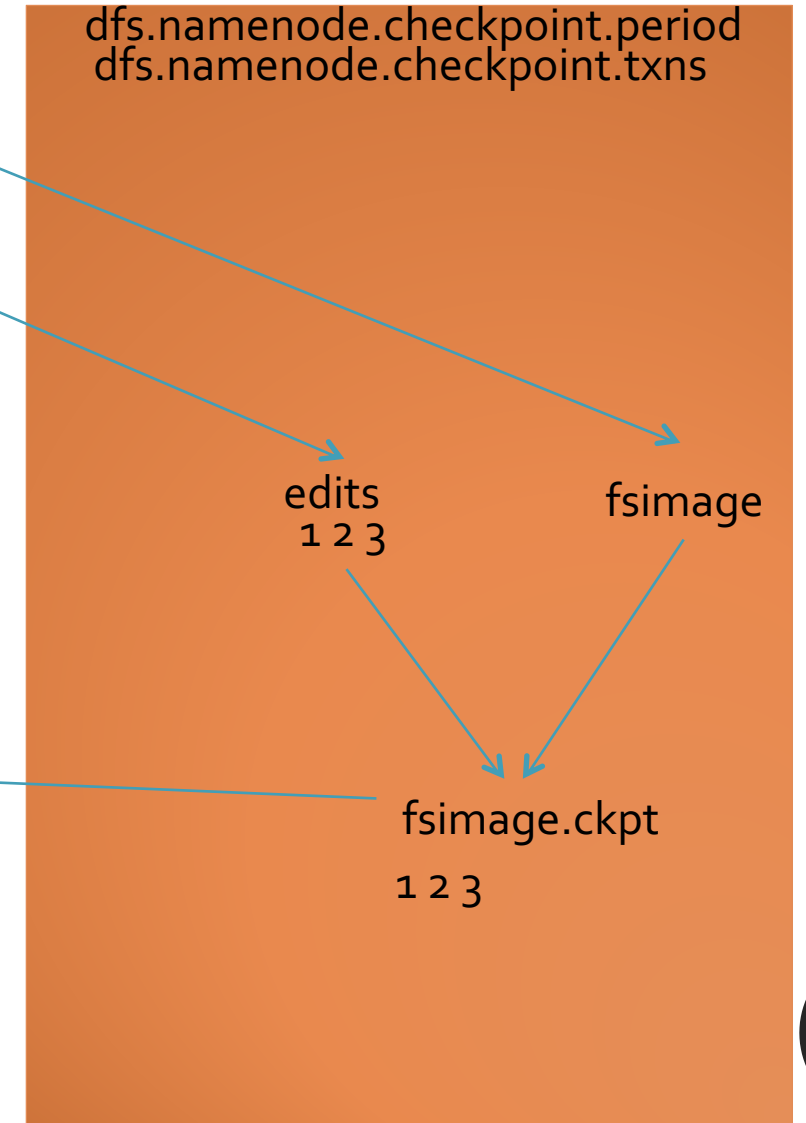


# Understanding Hadoop Cluster

## Name Node



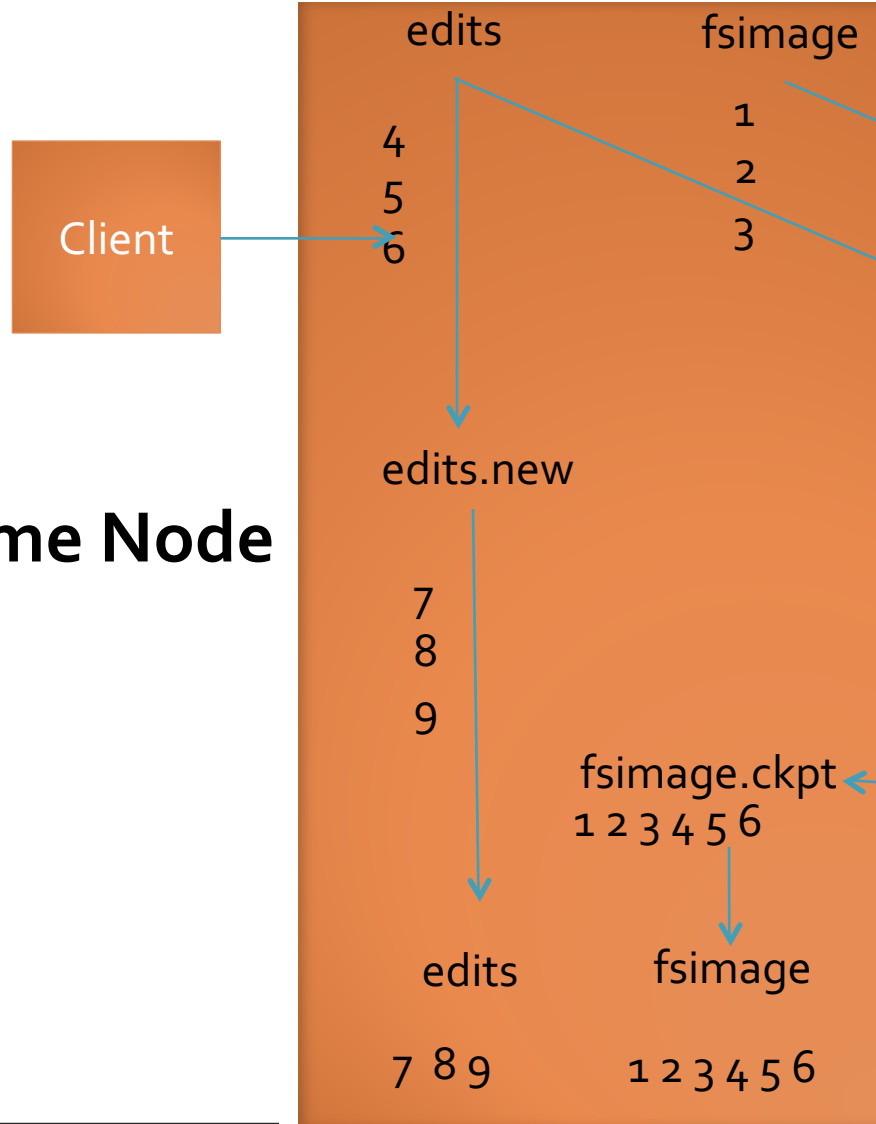
## Secondary Name Node






# Understanding Hadoop Cluster

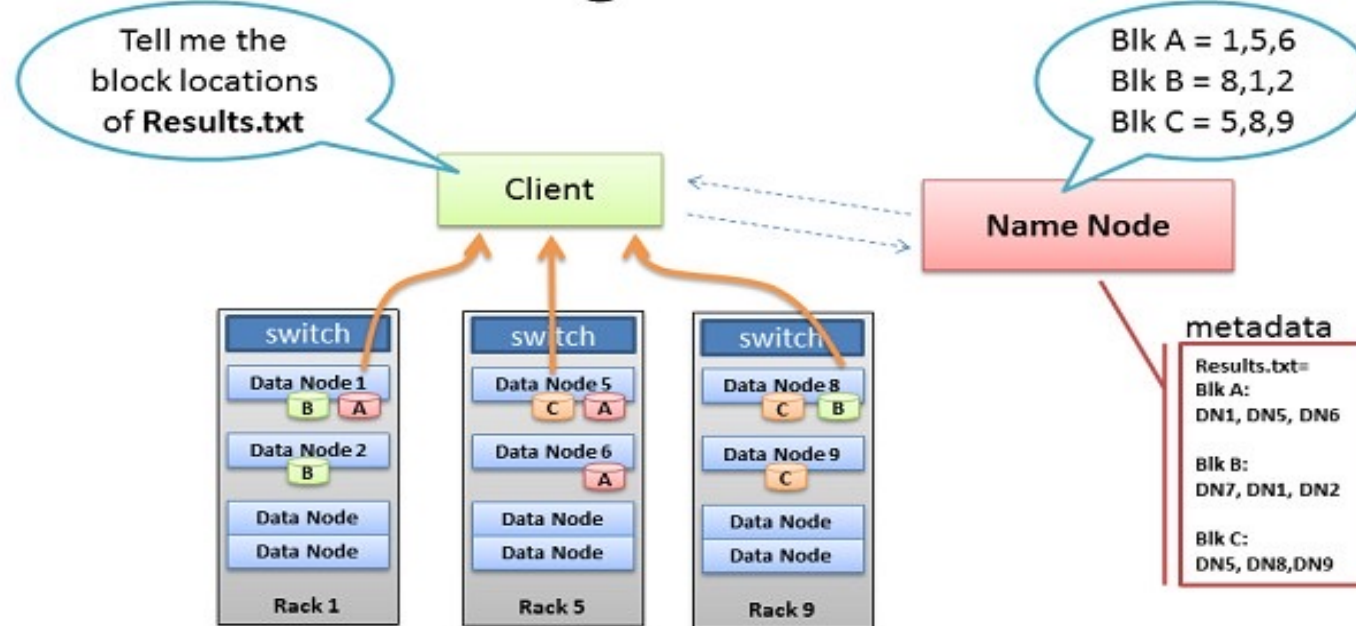
## Name Node



## Secondary Name Node



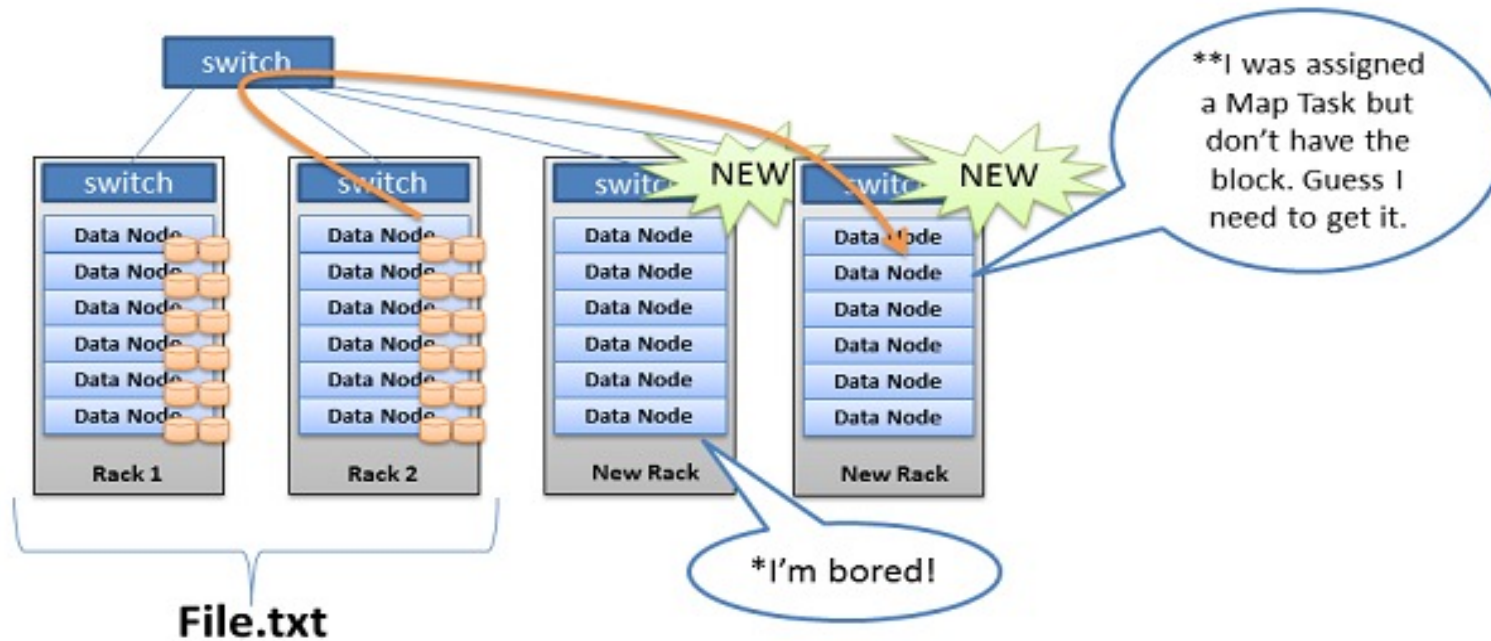
## Client reading files from HDFS



- Client receives Data Node list for each block
- Client picks first Data Node for each block
- Client reads blocks sequentially

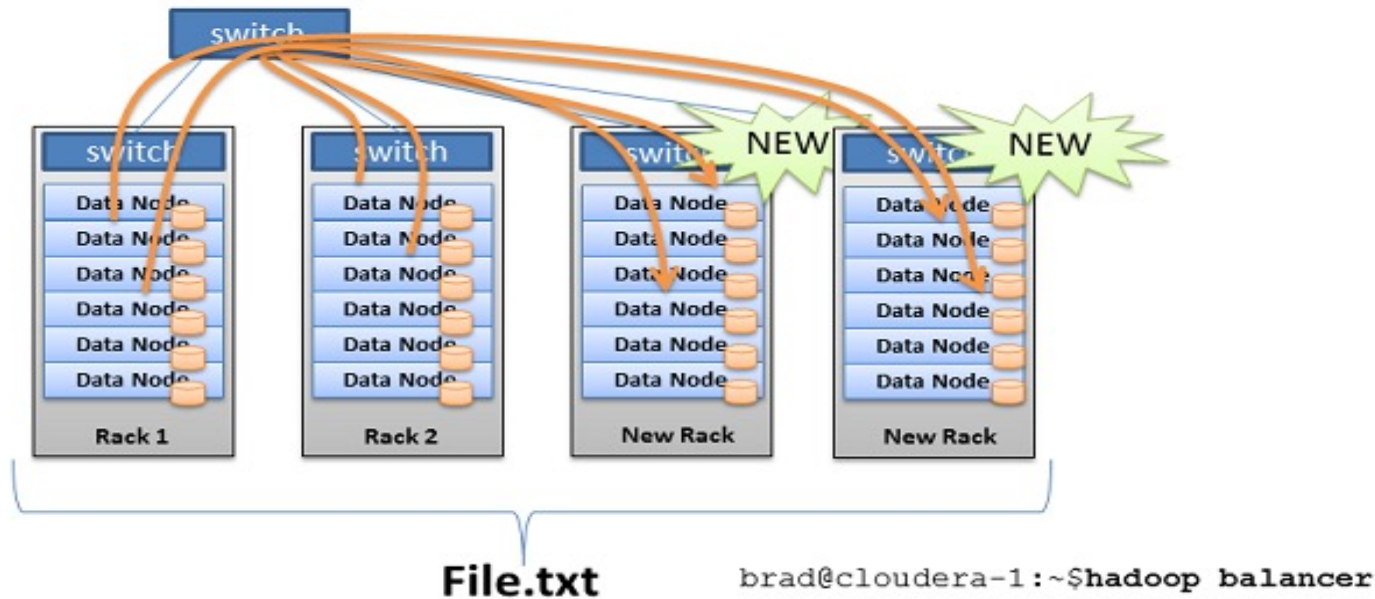


## Unbalanced Cluster



- Hadoop prefers local processing if possible
- New servers underutilized for Map Reduce, HDFS\*
- More network bandwidth, slower job times\*\*

## Cluster Balancing



- Balancer utility (if used) runs in the background
- Does not interfere with Map Reduce or HDFS
- Default rate limit 1 MB/s

# General Planning Considerations



Handle  
More  
Data



At  
Lower  
Cost



In  
Less  
Time



With  
Less  
Power

☁ Start with small cluster ( 4 to 10 nodes) and grow as and when required. Cluster can be grown whenever there is a

- Increase in computation power needed
- Increase in data to be stored
- Increase in amount of memory to process tasks
- Increase in data transfer between data nodes

Cluster Growth based on Storage Capacity:

Data Growth TB/Week	Replication Factor	Intermediate & Log Files	Overall Space needed per week
2	3	30%	7.8

Two Machines with 1X4TB are needed.

# Where to Optimize



## Hardware



## Software



# Choosing Right Hardware



## Master Node:

- ☁ Single Point of Failure
- ☁ 32 GB RAM
- ☁ Dual Xeon E5600 or better (Quad core)
- ☁ Dual Power supply for Redundancy
- ☁ 4 x 500 GB 7200 rpm SATA drives
- ☁ Dual 1 Gb Ethernet cards

## Data Nodes:

- ☁ 4 1TB hard disks in a JBOD (Just a Bunch Of Disks) configuration. No RAID.
- ☁ 2 quad core CPUs, running at least 2-2.5GHz
- ☁ 16-24GBs of RAM (24-32GBs if you're considering HBase)
- ☁ Gigabit Ethernet

- ☁ Master Node:
- ☁ No Commodity Hardware
- ☁ RAIDed hard drives
- ☁ Backup Metadata to an NFS Mount
- ☁ RAM Thumb rule: 1 GB per 1 million blocks of data. 32GB for 100 nodes.
- ☁ If Metadata is lost, whole cluster is lost. Use expensive Name Node.

- ☁ # of Tasks per Core:
- ☁ 2 Cores - Datanode and Tasktracker
- ☁ Thumb Rule – 1 Core can run 1.5 Mappers or Reducers
- ☁ Amount of RAM:
- ☁ Thumb Rule: 1G per Map or Reduce task

## Choosing Right Hardware Based on Different workloads



**Light Processing Configuration (1U/machine):** Two quad core CPUs, 8GB memory, and 4 disk drives (1TB or 2TB). Note that CPU-intensive work such as natural language processing involves loading large models into RAM before processing data and should be configured with 2GB RAM/core instead of 1GB RAM/core.


### Balanced Compute Configuration

**(1U/machine):** Two quad core CPUs, 16 to 24GB memory, and 4 disk drives (1TB or 2TB) directly attached using the motherboard controller. These are often available as twins with two motherboards and 8 drives in a single 2U cabinet.

**Storage Heavy Configuration (2U/machine):** Two quad core CPUs, 16 to 24GB memory, and 12 disk drives (1TB or 2TB). The power consumption for this type of machine starts around ~200W in idle state and can go as high as ~350W when active.

### Compute Intensive Configuration

**(2U/machine):** Two quad core CPUs, 48-72GB memory, and 8 disk drives (1TB or 2TB). These are often used when a combination of large in-memory models and heavy reference data caching is required.





## Hadoop HDFS Commands

---